

**COMPARE  
LOADBALANCING  
AUTOSCALING  
SERVERLESS COMPUTING**



**AWS**

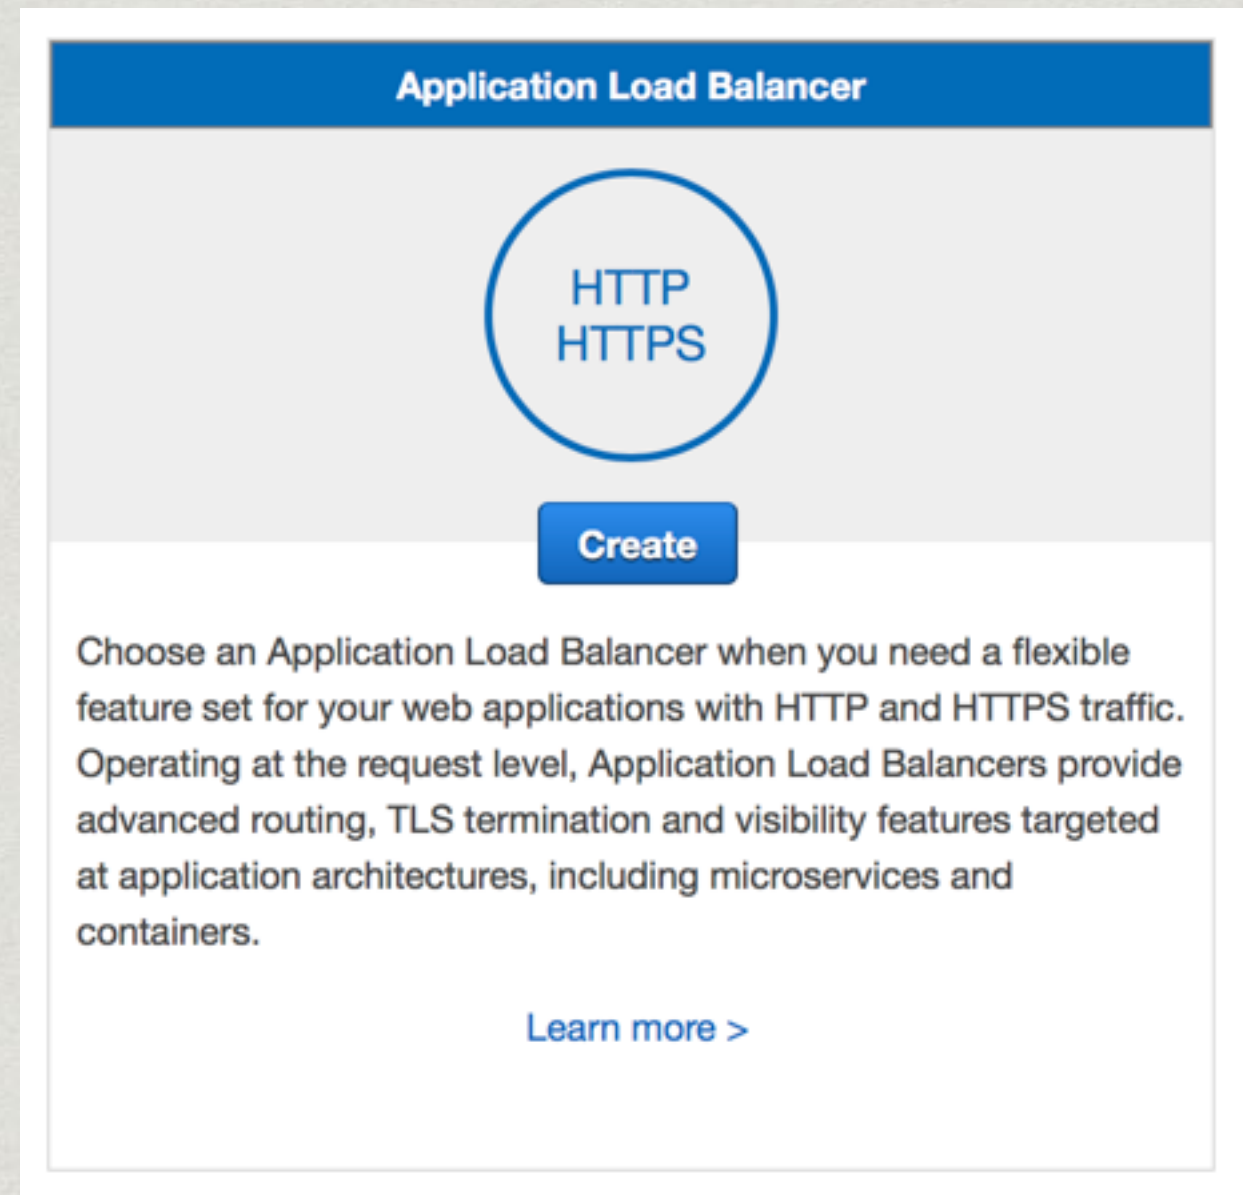
# AWS Load Balancing

- \* Elastic Load Balancing supports three types of load balancers:
- \* Application Load Balancers
- \* Network Load Balancers
- \* Classic Load Balancers



# AWS Application Load Balancers

- \* Application Load Balancer is best suited for load balancing of HTTP and HTTPS traffic and provides advanced request routing targeted at the delivery of modern application architectures, including microservices and containers. Operating at the individual request level (Layer 7), Application Load Balancer routes traffic to targets within Amazon Virtual Private Cloud (Amazon VPC) based on the content of the request.



The graphic features a blue header bar with the text "Application Load Balancer". Below this, a light gray rectangular area contains a blue circle with the text "HTTP" and "HTTPS" stacked vertically. Underneath the circle is a blue button with the word "Create" in white. Below the gray area, on a white background, is a paragraph of text describing the service's capabilities for web applications, microservices, and containers. At the bottom right of the white area is a blue link that says "Learn more >".

**Application Load Balancer**

HTTP  
HTTPS

**Create**


Choose an Application Load Balancer when you need a flexible feature set for your web applications with HTTP and HTTPS traffic. Operating at the request level, Application Load Balancers provide advanced routing, TLS termination and visibility features targeted at application architectures, including microservices and containers.

[Learn more >](#)

# AWS Network Load Balancer

- \* Network Load Balancer is best suited for load balancing of TCP traffic where extreme performance is required. Operating at the connection level (Layer 4), Network Load Balancer routes traffic to targets within Amazon Virtual Private Cloud (Amazon VPC) and is capable of handling millions of requests per second while maintaining ultra-low latencies. Network Load Balancer is also optimized to handle sudden and volatile traffic patterns.

## Network Load Balancer



Create

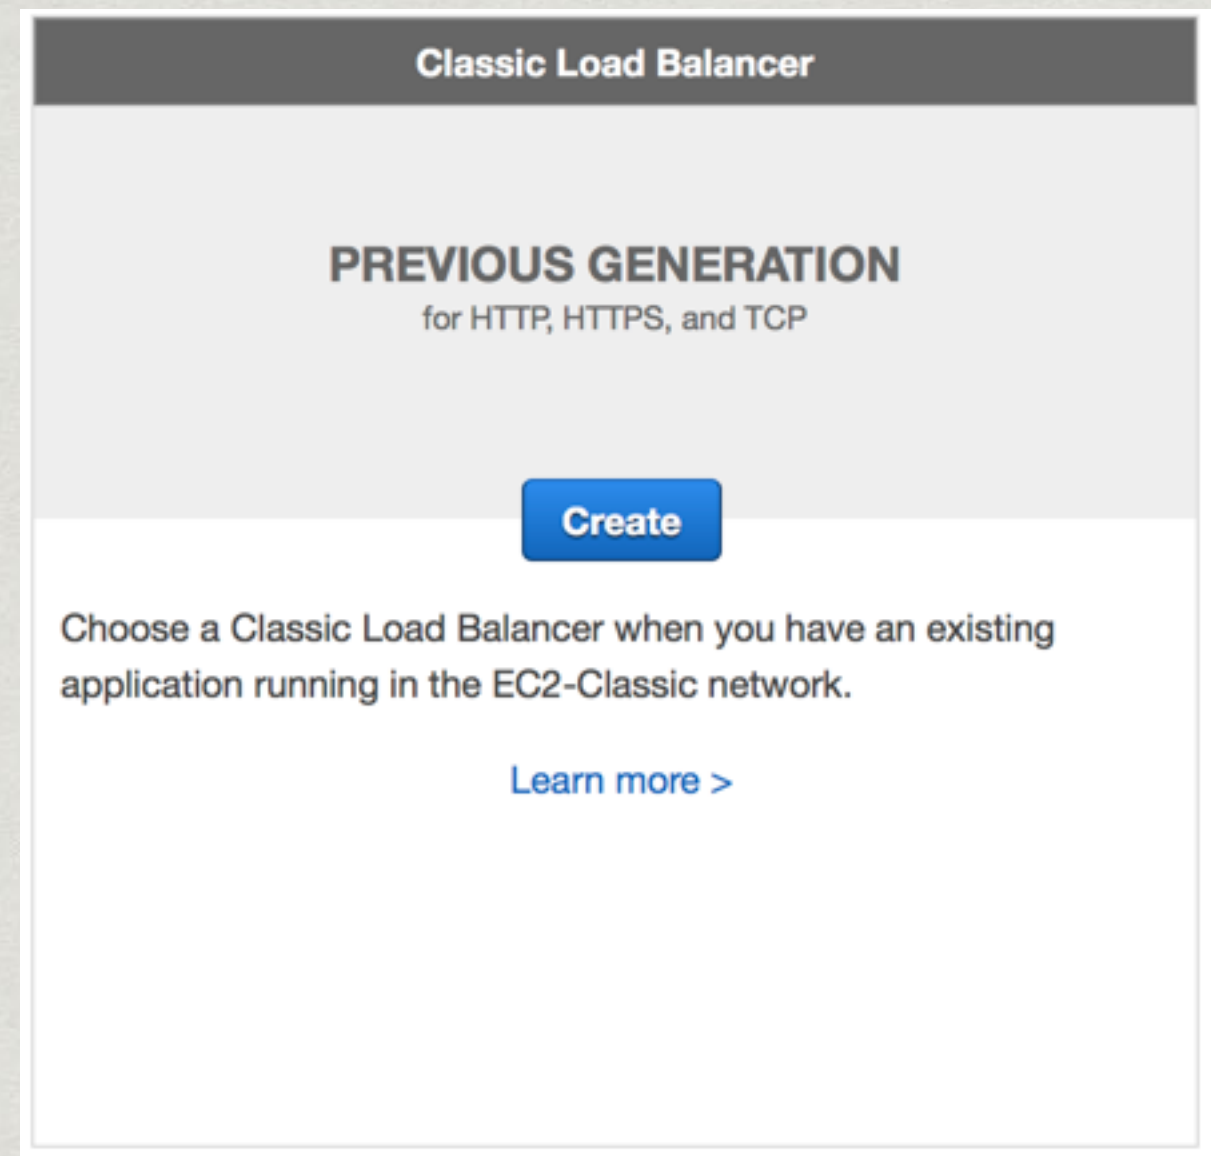
Choose a Network Load Balancer when you need ultra-high performance and static IP addresses for your application. Operating at the connection level, Network Load Balancers are capable of handling millions of requests per second while maintaining ultra-low latencies.

[Learn more >](#)



# AWS Classic Load Balancer

- \* Classic Load Balancer provides basic load balancing across multiple Amazon EC2 instances and operates at both the request level and connection level. Classic Load Balancer is intended for applications that were built within the EC2-Classic network.



# AWS Auto Scaling

- \* Amazon EC2 Auto Scaling is designed to automatically launch or terminate EC2 instances based on user-defined policies, schedules, and health checks. Use this service in conjunction with the AWS Auto Scaling, Amazon CloudWatch, and Elastic Load Balancing services.



# AWS Auto Scaling

aws

Services

Resource Groups

Admin/jdoe

N. Virginia

Support

Step 1

Select application

Step 2

Configure scaling plan

Step 3

Review and create

Application Auto Scaling > Create scaling plan

Select application

Select the AWS CloudFormation stack for your application. You cannot use a stack that is in a failed or deleted state.

CloudFormation stacks (3)

	Name	Status	Creation time
<input type="radio"/>	My-App-1	CREATE_COMPLETE	2017-11-02 19:57:34 UTC-0700
<input type="radio"/>	My-App-2	CREATE_COMPLETE	2017-12-11 04:12:01 UTC-0700
<input type="radio"/>	My-App-3	CREATE_COMPLETE	2018-01-02 10:57:09 UTC-0700

Cancel

Next



# AWS Serverless

- \* COMPUTE: AWS Lambda
- \* API PROXY: Amazon API Gateway
- \* STORAGE: Amazon Simple Storage Service (Amazon S3)
- \* DATA STORES: Amazon DynamoDB
- \* INTERPROCESS MESSAGING: Amazon SNS SQS
- \* ORCHESTRATION: AWS Step Functions
- \* ANALYTICS: Amazon Kinesis, Amazon Athena

# Azure



# Azure Load Balancer

- \* Load balance incoming Internet traffic to virtual machines. This configuration is known as a **public Load Balancer**.
- \* Load balance traffic between virtual machines inside a virtual network. You can also reach a Load Balancer frontend from an on-premises network in a hybrid scenario. Both of these scenarios use a configuration that is known as an **internal Load Balancer**.
- \* Port forward traffic to a specific port on specific virtual machines with inbound NAT rules.
- \* Provide **outbound connectivity** for virtual machines inside your virtual network by using a public Load Balancer.

# Autoscaling


- \* **Virtual Machines** support autoscaling through the use of [VM Scale Sets](#), which are a way to manage a set of Azure virtual machines as a group.
- \* **Service Fabric** also supports auto-scaling through VM Scale Sets. Every node type in a Service Fabric cluster is set up as a separate VM scale set. That way, each node type can be scaled in or out independently.
- \* **Azure App Service** has built-in autoscaling. Autoscale settings apply to all of the apps within an App Service.
- \* **Azure Cloud Services** has built-in autoscaling at the role level.
- \* **Azure Functions** differs from the previous compute options, because you don't need to configure any autoscale rules. Instead, Azure Functions automatically allocates compute power when your code is running, scaling out as necessary to handle load.



# Azure Serverless Computing

# Timer-based processing

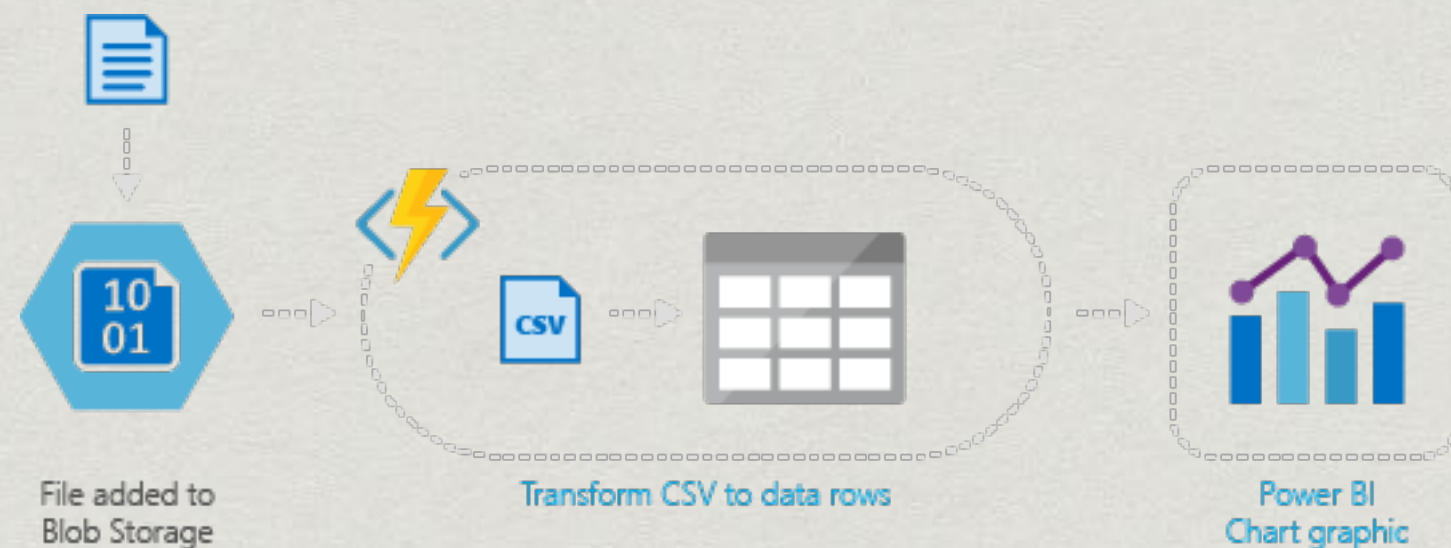
- \* Functions supports an event based on a timer using Cron job syntax. For example, execute code that runs every 15 minutes and clean up a database table based on custom business logic.





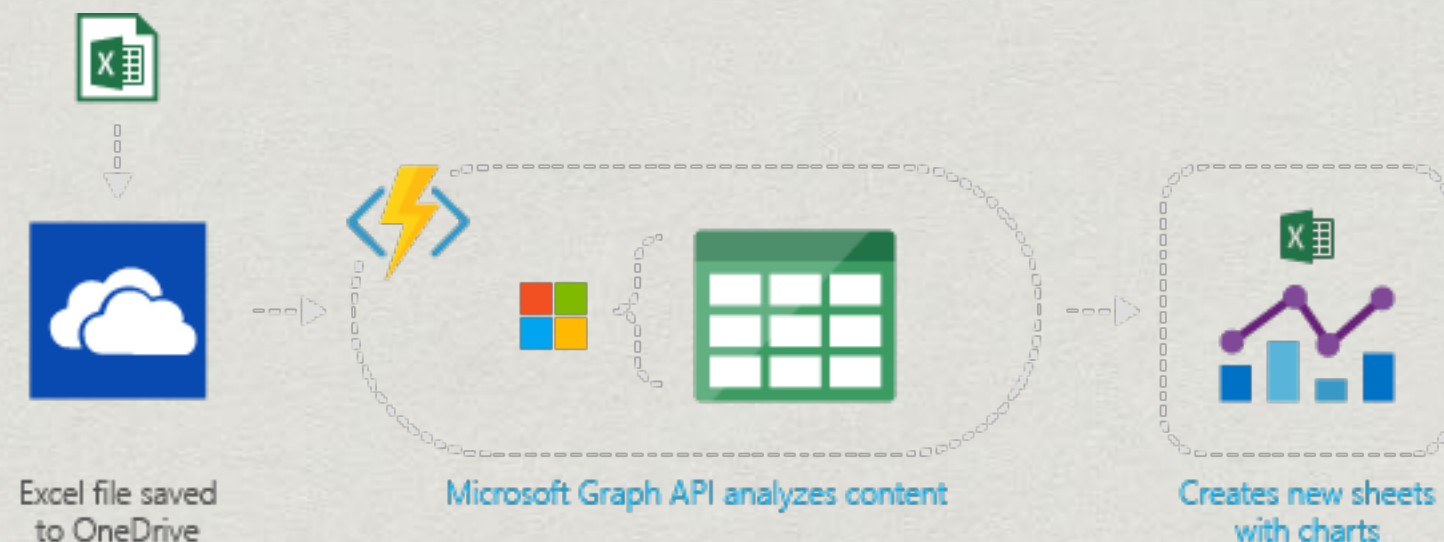
# Azure service event processing

- \* Functions supports triggering an event based on an activity in an Azure service. For example, execute serverless code that reads newly discovered test log files in a Blob storage container, and transform this into a row in a SQL Database table.



# SaaS event processing


- \* Functions supports triggers based on activity in a Software as a service (SaaS)-based application. For example, save a file in OneDrive, which triggers a function that uses the Microsoft Graph API to modify the spreadsheet, and creates additional charts and calculated data.





# Serverless web application architectures

- \* Functions can power a single-page app. The app calls functions using the WebHook URL, saves user data, and decides what data to display. Or, do simple customizations, such as changing ad targeting by calling a function and passing it user profile information.



# Serverless mobile back ends

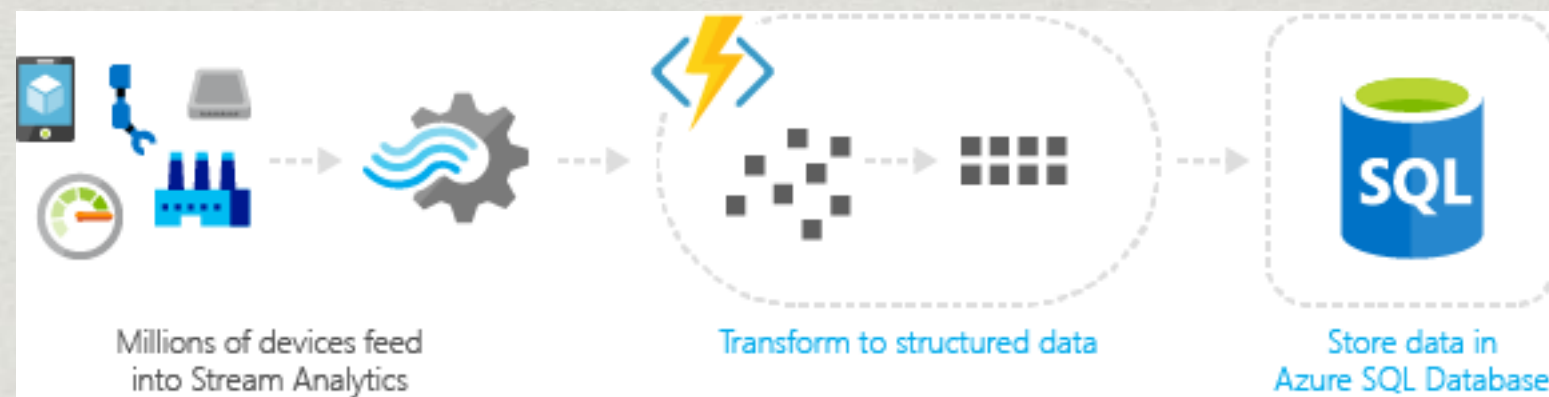
- \* A mobile back end can be a set of HTTP APIs that are called from a mobile client using the WebHook URL. For example, a mobile application can capture an image, and then call a function to get an access token for uploading to blob storage. A second function is triggered by the blob upload and resizes the image to be mobile-friendly.





# Real-time stream processing

- \* For example, Internet of Things devices send messages to Stream Analytics, which then calls a function to transform the message. This function processes the data and creates a new record in a SQL database.



# Real-time bot messaging

- \* Use Functions to customize the behavior of a bot using a WebHook. For example, create a function that processes a message using Cortana Analytics and call this function using Microsoft Bot Framework.





# Google Cloud Compute

# Global external load balancing

- \* **HTTP(S) load balancing** distributes HTTP(S) traffic among groups of instances based on proximity to the user, the requested URL, or both.
- \* **SSL Proxy load balancing** distributes SSL traffic among groups of instances based on proximity to the user.
- \* **TCP Proxy load balancing** distributes TCP traffic among groups of instances based on proximity to the user.



# Regional external load balancing

- \* **Network load balancing** distributes traffic among a pool of instances within a region. Network load balancing can balance any kind of TCP/UDP traffic.

# Regional internal load balancing

- \* **Internal load balancing** distributes traffic from Google Cloud Platform virtual machine instances to a group of instances in the same region.



# Autoscaling

- \* CPU utilization: CPU utilization is the most basic autoscaling that you can perform. This policy tells the autoscaler to watch the average CPU utilization of a group of virtual machines and add or remove virtual machines from the group to maintain your desired utilization. This is useful for configurations that are CPU-intensive but might fluctuate in CPU usage.
- \* Load balancing serving capacity: Set up an autoscaler to scale based on load balancing serving capacity and the autoscaler will watch the serving capacity of an instance group, and scale if the virtual machines are over or under capacity. The serving capacity of an instance can be defined in the load balancer's [backend service](#) and can be based on either utilization or requests per second.
- \* Stackdriver Monitoring metrics: If you export or use Stackdriver Monitoring metrics, you can set up autoscaling to collect data of a specific metric and perform scaling based on your desired utilization level. It is possible to scale based on standard metrics provided by Stackdriver Monitoring, or using any custom metrics you create as well.

# Serverless application backends

Trigger your code from GCP services or call it directly from any web, mobile, or backend application



## Integration with third-party services and APIs

Surface your own microservices or quickly extend your application by integrating with third-party services

[LEARN MORE](#)



## Serverless mobile backends

Extend your Firebase application functionality with powerful compute, data analytics, and machine learning capabilities without spinning up a server

[LEARN MORE](#)



## Serverless IoT backends

Build backends for Internet of Things (IoT) device telemetry data collection, real-time processing, and analysis

[LEARN MORE](#)



# Real-time data processing systems

Run your code in response to changes in data



## Real-time file processing

Execute your code in response to changes in data in Cloud Storage buckets

[LEARN MORE](#)



## Real-time stream processing

Respond to events from Cloud Pub/Sub to process, transform, and enrich streaming data

[LEARN MORE](#)



## Event-driven extract, transform, load (ETL)

Retrieve and pre-process data and load transformed datasets for further querying and analysis

[LEARN MORE](#)

# Intelligent applications

Easily inject artificial intelligence into your applications



## Virtual assistants and chatbots

Extend your products and services with voice and text-based natural conversational experiences that help users get things done

[LEARN MORE](#)



## Video and image analysis

Retrieve relevant information from videos and images to search, discover, and derive insight from your media content

[LEARN MORE](#)



## Sentiment analysis

Reveal the structure and meaning of text and add sentiment analysis and intent extraction capabilities to your applications

[LEARN MORE](#)